

LAB 10 - Further topics

Iulia Cioroianu

24/07/2020

1. Reading in data

```
require(quanteda)
```

```
## Loading required package: quanteda
## Package version: 2.1.0
## Parallel computing: 2 of 16 threads used.
## See https://quanteda.io for tutorials and examples.
##
## Attaching package: 'quanteda'
## The following object is masked from 'package:utils':
##
##      View
```

```
require(quanteda.textmodels)
```

```
## Loading required package: quanteda.textmodels
##
## Attaching package: 'quanteda.textmodels'
## The following object is masked from 'package:quanteda':
##
##      data_dfm_lbgexample
```

```
require(topicmodels)
```

```
## Loading required package: topicmodels
```

```
require(stm)
```

```
## Loading required package: stm
## stm v1.3.5 successfully loaded. See ?stm for help.
## Papers, resources, and other materials at structuraltopicmodel.com
```

```
require(lubridate)
```

```
## Loading required package: lubridate
##
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library(corrplot)

## corrplot 0.84 loaded
articles <- read.csv("diffbot_text_basic.csv", stringsAsFactors = FALSE, encoding = "utf-8")

colnames(articles) <- c("X", "Date", "Sentiment", "Region", "Country", "Site", "Type", "Title", "text")

# In the first part of the analysis we will focus on the UK.
art_uk <- subset(articles, Country == "United Kingdom")

# Change date to simpler format
art_uk$Date2 <- dmy(substr(art_uk$Date, 6, 16))

# Define the year and discard all articles before 2020
art_uk$year <- year(art_uk$Date2)
art_uk <- subset(art_uk, year >= 2020)

# Redefine year, month, week
art_uk$year <- year(art_uk$Date2)
art_uk$month <- month(art_uk$Date2)
art_uk$week <- week(art_uk$Date2)

# Create our corpus
corp_uk <- corpus(art_uk)
```

2. Processing: Include compounding multi-word expressions based on collocation analysis

```
# Tokenize corpus
toks <- tokens(corp_uk,
               remove_numbers=TRUE,
               remove_punct=TRUE,
               remove_symbols=TRUE,
               remove_separators=TRUE,
               remove_url = TRUE,
               verbose=TRUE)
```

```
## Creating a tokens object from a corpus input...
## ...starting tokenization
## ...text1 to text666
## ...preserving hyphens
## ...preserving social media tags (#, @)
## ...segmenting into words
## ...22,709 unique types
## ...removing separators, punctuation, symbols, numbers, URLs
## ...complete, elapsed time: 1.36 seconds.
## Finished constructing tokens from 666 documents.
# Remove stopwords
toks <- tokens_remove(toks, pattern = stopwords('en'))

# Identify collocations
tstat_col_cap <- textstat_collocations(toks, size=2, min_count = 10, tolower = FALSE)

head(tstat_col_cap, 20)
```

	collocation	count	count_nested	length	lambda	z
## 1	contact tracing	1458	0	2	6.421352	115.61147
## 2	Apple Google	575	0	2	6.600089	89.11548
## 3	public health	389	0	2	4.989355	73.36020
## 4	tested positive	208	0	2	6.562935	62.88122
## 5	human rights	185	0	2	7.665732	59.31179
## 6	Mr Hancock	179	0	2	6.210423	58.66975
## 7	Google Apple	234	0	2	4.737815	58.48396
## 8	track trace	161	0	2	6.130155	55.97554
## 9	contact-tracing app	396	0	2	4.139739	54.90559
## 10	around world	145	0	2	6.111660	54.19148
## 11	social distancing	277	0	2	8.855249	54.06475
## 12	location data	279	0	2	4.523757	54.05236
## 13	come contact	246	0	2	4.551813	53.38716
## 14	tracing app	490	0	2	2.827168	53.08473
## 15	Secretary Matt	118	0	2	7.933278	50.87028
## 16	Health Secretary	175	0	2	7.441061	49.64031
## 17	make sure	136	0	2	6.961924	49.15818
## 18	mobile phone	143	0	2	5.184597	49.11494
## 19	tracing apps	224	0	2	3.781514	48.97105
## 20	spread virus	156	0	2	4.667330	48.67244

Compound multi-word expressions

Results of collocation analysis can be use it to compound tokens. We will compound strongly associated multi-word expressions by sub-setting `tstat_col_cap$collocation`.

Collocations are automatically recognized as multi-word expressions by `tokens_compound()` in case-sensitive fixed pattern matching. This is the fastest way to compound large numbers of multi-word expressions, but make sure that `tolower = FALSE` in `textstat_collocations()` to do this.

```
toks_comp <- tokens_compound(toks, pattern = tstat_col_cap[tstat_col_cap$z > 3])
# Text1 without compounding
toks[['text1']][1:50]
```

```
## [1] "NHS"           "begun"          "feeding"         "health"
## [5] "workers"       "use"            "personal"        "protective"
## [9] "equipment"     "PPE"           "data"            "store"
## [13] "system"        "designed"       "identify"        "hospitals"
## [17] "GP"            "surgeries"      "risk"            "running"
## [21] "kit"           "address"        "problem"         "occurs"
## [25] "High-level"    "decision-makers" "able"            "start"
## [29] "seeing"        "information"     "via"             "computer"
## [33] "dashboard"     "within"         "fortnight"       "NHS"
## [37] "staff"         "say"            "lives"           "put"
## [41] "risk"          "PPE"            "shortages"       "government"
## [45] "said"          "working"        "around"          "clock"
## [49] "address"       "issue"
```

```
# Text 1 with compounding
toks_comp[['text1']][1:50]
```

```
## [1] "NHS"           "begun"
## [3] "feeding"       "health_workers"
## [5] "use_personal_protective_equipment" "PPE"
## [7] "data_store"    "system"
## [9] "designed"      "identify"
## [11] "hospitals"     "GP"
## [13] "surgeries"     "risk"
## [15] "running"       "kit"
## [17] "address"       "problem"
## [19] "occurs"        "High-level"
## [21] "decision-makers" "able"
## [23] "start"         "seeing"
## [25] "information"   "via"
## [27] "computer"      "dashboard"
## [29] "within"        "fortnight"
## [31] "NHS_staff"     "say"
## [33] "lives"         "put_risk"
## [35] "PPE"           "shortages"
## [37] "government_said_working" "around"
## [39] "clock"         "address"
## [41] "issue"         "NHS"
## [43] "Providers"     "represents"
## [45] "hospitals"     "NHS"
## [47] "trusts"        "England"
## [49] "told_BBC"      "supplies"
```

3. Describing the data: Finding words associated with a certain word.

We can find words associated with target words using the window argument of `tokens_select()`.

```

# Create subset of tokens around the term "privacy"
toks_privacy <- tokens_keep(toks_comp, pattern = 'privacy', window = 10) # equivalent to tokens_select(
# Create subset of tokens which are not around the term "privacy"
toks_noprivacy <- tokens_remove(toks_comp, pattern = 'privacy', window = 10) # equivalent to tokens_sel

# Turn both of them into a DFM
dfmat_privacy <- dfm(toks_privacy)
dfmat_noprivacy <- dfm(toks_noprivacy)

# Calculate keyness between the two categories - the score for features that occur differentially
# across the two categories
tstat_key_privacy <- textstat_keyness(rbind(dfmat_privacy, dfmat_noprivacy), seq_len(ndoc(dfmat_privacy)
tstat_key_privacy_subset <- tstat_key_privacy[tstat_key_privacy$n_target > 10, ]
head(tstat_key_privacy_subset, 50)

```

##	feature	chi2	p	n_target	n_reference
## 1	privacy	14582.99461	0.000000e+00	734	0
## 3	preserving	154.55254	0.000000e+00	13	6
## 4	covid-19_exposure_logging	129.44714	0.000000e+00	11	5
## 16	assurances	88.92404	0.000000e+00	19	41
## 22	default	73.76737	0.000000e+00	12	18
## 23	settings	73.16860	0.000000e+00	21	61
## 25	thinks	67.96699	1.110223e-16	13	24
## 31	information_commissioner	62.27760	2.997602e-15	15	36
## 43	invasive	58.88863	1.665335e-14	17	49
## 51	job	48.38537	3.501754e-12	21	90
## 52	debate	46.83615	7.717715e-12	15	47
## 59	civil_liberties	45.70936	1.371658e-11	15	48
## 87	design	39.35906	3.526137e-10	21	104
## 94	controversial	37.47155	9.275609e-10	11	31
## 95	assessment	37.47155	9.275609e-10	11	31
## 96	concern	36.53211	1.501718e-09	18	80
## 98	balance	36.18691	1.792706e-09	11	32
## 99	amnesty_international	36.18691	1.792706e-09	11	32
## 106	legal	32.44274	1.227569e-08	19	100
## 114	protecting	29.59496	5.324320e-08	13	52
## 133	governments	28.15931	1.117289e-07	33	254
## 137	designed	27.77797	1.360670e-07	17	92
## 138	concerned	27.77797	1.360670e-07	17	92
## 139	questions	27.42790	1.630614e-07	25	170
## 140	human_rights	26.96172	2.075247e-07	23	151
## 144	freedom	26.33593	2.869017e-07	12	49
## 147	comment	25.31931	4.858208e-07	11	43
## 148	contact_tracing_apps	25.18626	5.205161e-07	18	107
## 231	fundamental	21.15676	4.232052e-06	11	49
## 233	ethical	20.54840	5.814232e-06	11	50
## 234	privacy_security	20.54840	5.814232e-06	11	50
## 243	mean	19.38133	1.070485e-05	15	88
## 246	choice	18.84684	1.416457e-05	11	53
## 250	analysis	18.43820	1.755044e-05	16	106
## 269	location_data	17.26484	3.251487e-05	18	131
## 270	expert	17.25818	3.262900e-05	14	84
## 271	citizens	16.11703	5.954605e-05	20	158
## 286	chair	15.06419	1.039163e-04	11	61

## 290	comes	14.88770	1.141045e-04	19	152
## 291	tech	14.82012	1.182667e-04	23	200
## 299	risk	14.58678	1.338503e-04	40	426
## 301	groups	14.43103	1.453872e-04	15	109
## 302	concerns	14.27970	1.575552e-04	25	228
## 303	apple	14.01793	1.810761e-04	24	217
## 307	app	13.53313	2.343883e-04	133	1900
## 309	believe	13.44512	2.456445e-04	16	124
## 352	needs	12.75175	3.556764e-04	21	187
## 357	often	12.44982	4.180310e-04	11	68
## 366	systems	11.86525	5.719103e-04	17	143
## 369	law	11.43750	7.197649e-04	21	195

These are the terms that are more likely to occur around the term “privacy” than anywhere else in the text.

4. Dictionary methods: Targeted sentiment analysis

You can use `tokens_select()` with `window` argument to perform *TARGETED* sentiment analysis.

Let’s evaluate the sentiment around mentions of privacy and security.

```
# Define keyterms
privacy <- c('privacy', 'security')

# Define relevant tokens - 20 tokens before and after the keyword
toks_privacy <- tokens_keep(toks_comp, pattern = phrase(privacy), window = 20)
toks_privacy
```

```
## Tokens consisting of 666 documents and 12 docvars.
## text1 :
## [1] "involved" "NHSX's"
## [3] "coronavirus_contact-tracing_app" "two"
## [5] "efforts" "otherwise"
## [7] "independent" "plans"
## [9] "mix" "information_gathered"
## [11] "via_app" "data_store"
## [ ... and 70 more ]
##
## text2 :
## character(0)
##
## text3 :
## character(0)
##
## text4 :
## [1] "unprecedented" "collaborations" "government" "tech_giants"
## [5] "sounds" "familiar" "pre-Covid" "precise"
## [9] "app-driven" "gig-fuelled" "future" "sold"
## [ ... and 170 more ]
##
## text5 :
## [1] "said" "see" "think"
## [4] "inevitable" "view" "crucially"
## [7] "planning" "around" "want"
```

```
## [10] "position"                  "protective_equipment" "face_masks"
## [ ... and 29 more ]
##
## text6 :
## character(0)
##
## [ reached max_ndoc ... 660 more documents ]
# Put it into a dataframe matrix, defining the week as the grouping variable
dfmat_privacy_lsd <- dfm(toks_privacy, dictionary = data_dictionary_LSD2015[1:2]) %>%
  dfm_group(group = 'week', fill = TRUE)
dfmat_privacy_lsd
```


```
## Document-feature matrix of: 20 documents, 2 features (2.5% sparse) and 4 docvars.
##      features
## docs negative positive
##  6           0         4
##  9           3         8
## 10          10        19
## 11          35        32
## 12          48        45
## 13          78       126
## [ reached max_ndoc ... 14 more documents ]
```

Let's plot the results. Start with positive and negative sentiment.

```
# What is the minimum date?
min(art_uk$Date2)
```


```
## [1] "2020-02-11"
```

```
matplot(dfmat_privacy_lsd, type = 'l', xaxt = 'n', lty = 1, ylab = 'Frequency')
grid()
axis(1, seq_len(ndoc(dfmat_privacy_lsd)), ymd("2020-02-11") + weeks(seq_len(ndoc(dfmat_privacy_lsd)) - 1))
legend('topleft', col = 1:2, legend = c('Negative', 'Positive'), lty = 1, bg = 'white')
```



Now plot relative sentiment:

```
n_eu <- ntoken(dfm(toks_privacy, group = toks_privacy$week))
plot((dfmat_privacy_lsd[,2] - dfmat_privacy_lsd[,1]) / n_eu,
     type = 'l', ylab = 'Sentiment', xlab = '', xaxt = 'n')
axis(1, seq_len(ndoc(dfmat_privacy_lsd)), ymd("2020-02-11") + weeks(seq_len(ndoc(dfmat_privacy_lsd)) - 1))
grid()
abline(h = 0, lty = 2)
```




Exercise:

Work with your group to measure and plot the use of terms related to China and Wuhan in a window of 15 tokens before and after the keywords in this corpus. #####

```
# Define the dictionary
gov <- c('china', 'China', 'chinese', 'Wuhan')
toks_gov <- tokens_keep(toks_comp, pattern = phrase(gov), window = 20)
dfmat_gov_lsd <- dfm(toks_gov, dictionary = data_dictionary_LSD2015[1:2]) %>%
  dfm_group(group = 'week', fill = TRUE)

matplot(dfmat_gov_lsd, type = 'l', xaxt = 'n', lty = 1, ylab = 'Frequency')
grid()
axis(1, seq_len(ndoc(dfmat_gov_lsd)), ymd("2020-01-11") + weeks(seq_len(ndoc(dfmat_gov_lsd)) - 1))
legend('topleft', col = 1:2, legend = c('Negative', 'Positive'), lty = 1, bg = 'white')
```

```
n_gov <- ntoken(dfm(toks_gov, group = toks_gov$week))
plot((dfmat_gov_lsd[,2] - dfmat_gov_lsd[,1]) / n_gov,
     type = 'l', ylab = 'Sentiment', xlab = '', xaxt = 'n')
axis(1, seq_len(ndoc(dfmat_gov_lsd)), ymd("2020-02-11") + weeks(seq_len(ndoc(dfmat_gov_lsd)) - 1))
grid()
abline(h = 0, lty = 2)
```



5. Dictionary methods: Which words contribute towards dictionary category counts?

```
toks_new <- tokens(head(corp_uk, 3))
dfm_list <- list()
```

```

for (key in names(data_dictionary_LSD2015)) {
  this_dfm <- tokens_select(toks_new, data_dictionary_LSD2015[key], pad = TRUE) %>%
    tokens_compound(data_dictionary_LSD2015[key]) %>%
    tokens_replace("", "OTHER") %>%
    dfm(tolower = FALSE)
  dfm_list <- c(dfm_list, this_dfm)
}
names(dfm_list) <- names(data_dictionary_LSD2015)

dfm_list

## $negative
## Document-feature matrix of: 3 documents, 21 features (49.2% sparse) and 12 docvars.
##           features
## docs  risk problem shortages critical Dangerous warned dangerously alarm
## text1      2      2      1      1      1      1      1      1
## text2      0      0      0      0      0      1      0      0
## text3      0      0      0      0      0      1      0      0
##           features
## docs  Emergencies concern
## text1      1      1
## text2      0      0
## text3      0      0
## [ reached max_nfeat ... 11 more features ]
##
## $positive
## Document-feature matrix of: 3 documents, 40 features (55.0% sparse) and 12 docvars.
##           features
## docs  protective trusts help care innovation sense resources partner
## text1      1      2      4      3      1      1      1      1
## text2      0      0      2      0      0      0      0      0
## text3      0      0      2      0      0      0      0      0
##           features
## docs  protectors efforts
## text1      1      2
## text2      0      0
## text3      0      0
## [ reached max_nfeat ... 30 more features ]
##
## $neg_positive
## Document-feature matrix of: 3 documents, 1 feature (0.0% sparse) and 12 docvars.
##           features
## docs  OTHER
## text1  789
## text2  574
## text3  586
##
## $neg_negative
## Document-feature matrix of: 3 documents, 1 feature (0.0% sparse) and 12 docvars.
##           features
## docs  OTHER
## text1  789
## text2  574
## text3  586

```

7. Structural topic models

This method will only work on your own computer because of limited computer power provided by RStudio Cloud.

See the attached script “STM_example.R”.